

01 · Project Documentation

SDJ Editorial Portal — Project Documentation

Document	01 — Consolidated project documentation (main submission document)
Author	Roger Koranteng Obeng, student ID 22424140
Assessor	Prof. Solomon Mensah
Course	Advanced Software Engineering — individual final project
Date	2026-08-12
Live system	https://ugjcs-frontend.vercel.app · API https://tsxsbf9rzp.us-east-1.amazonaws.com
Repository	github.com/rogerkorantenng/ugjcs
Submitted as	Project_Documentation.pdf

Naming. The product documented here is the **SDJ Editorial Portal**, built for the **Science and Development Journal (SDJ)**, an established journal published by the College of Basic and Applied Sciences (CBAS), University of Ghana. The repository (github.com/rogerkorantenng/ugjcs), hosting URLs, package names and infrastructure resource names predate this naming and retain the pilot's internal codename **UGJCS**; wherever `ugjcs` appears in a path, URL or identifier in this document set, it is that codename, not a separate system. The deployed system is a prototype built for an Advanced Software Engineering exam — it is not SDJ's official production system.

Status of this document. This is a consolidating document. It summarises and cross-references six authoritative supporting documents and the codebase itself rather than duplicating them; where a figure, table, or claim has a fuller treatment elsewhere, that document is named and the reader is pointed to it. Nothing here should be taken to override what those documents say — where a discrepancy exists between two of them, it is recorded rather than silently resolved in one direction (see section 10.4).

Referenced source documents

#	Document	Role
02	docs/02-srs.md	Software Requirements Specification, traceability matrix
03	docs/03-effort-estimation.md	Use Case Points, COCOMO II cross-check, MoSCoW cut
04	docs/04-technical-debt-register.md	Fourteen debt entries, Debt→Cause→Impact→Priority→Resolution
05	docs/05-api-contract.md	HTTP boundary contract between backend and frontend
06	docs/06-testing-report.md	Test strategy, test cases, defects found, UAT
—	docs/superpowers/specs/2026-08-12-ugjcs-journal-platform-design.md	Design specification (problem, architecture, data model)
—	backend/src/ugjcs/domain/	The implemented domain — source of truth for lifecycle and rules
—	infra/	Terraform — source of truth for the deployed topology

1. Project title

SDJ Editorial Portal — a double-blind peer-review and editorial management platform built for the Science and Development Journal (SDJ), published by the College of Basic and Applied Sciences (CBAS), University of Ghana, as an individual Advanced Software Engineering final project within a 48-hour development window.

2. Problem statement

The College of Basic and Applied Sciences (CBAS), University of Ghana publishes the Science and Development Journal (SDJ), but the journal has no dedicated system for managing its editorial process: submissions, reviews and decisions are coordinated manually over email, shared drives and spreadsheets. That approach fails in four specific ways (design specification section 1):

1. **Blinding is not enforceable.** Double-blind review depends on a human remembering to strip identifying information from a document before forwarding it. Author names routinely survive in PDF metadata even when removed from the visible text.
2. **There is no audit trail.** When a rejected author appeals, there is no authoritative, tamper-evident record of who decided what, when, and on what evidence.
3. **Reviewer assignment is ad hoc.** Editors assign from memory, which concentrates load on a few willing reviewers and misses expertise matches entirely.
4. **Published work is not discoverable.** Accepted papers end up as files on a shared drive rather than in a citable, indexable, harvestable archive.

The portal addresses all four as first-class system responsibilities — enforced by the system — rather than as procedural guidance that depends on a human remembering to follow it.

3. Aim and objectives

Aim. Deliver a deployed, production-quality editorial portal for SDJ that manages the complete scholarly publishing lifecycle from submission to public archival, with double-blind integrity and editorial auditability enforced by the system rather than by convention.

Objectives:

- **O1.** Enforce a guarded manuscript lifecycle in which no illegal state transition is reachable through any interface.
- **O2.** Guarantee double-blind integrity structurally, including at the document level.
- **O3.** Provide a tamper-evident editorial audit trail.
- **O4.** Assist reviewer assignment with expertise matching, conflict-of-interest exclusion and workload balancing, leaving final authority with the editor.
- **O5.** Publish accepted work to a public, searchable, citable and machine-harvestable archive.
- **O6.** Demonstrate disciplined engineering practice: estimation-driven scope, automated quality gates, infrastructure as code, and an explicit technical debt register.

Objectives O1–O3 are demonstrably met by the delivered domain layer (section 8, section 10); O4 is partially met (reviewer assignment is a persistence record without matching or a conflict check — TD-02, TD-03); O5 is met for the archive read path; O6 is the subject of section 7, section 11 and section 12, and is, on the evidence in section 11.7, the objective this project delivers most convincingly.

4. Stakeholders

The client is the editorial office of the Science and Development Journal at CBAS; the primary human actors are the journal's editor-in-chief, its handling editors, its reviewers and its authors.

Actor	Type	Primary concerns
Author	Primary, authenticated	Submit work, track progress, respond to reviews
Reviewer	Primary, authenticated	Accept/decline invitations, submit structured reviews, remain anonymous
Editor	Primary, authenticated	Screen submissions, assign reviewers, record decisions
Editor-in-Chief	Primary, authenticated	Compose issues, publish, configure journal policy, final authority
Administrator	Secondary, authenticated	User accounts, role assignment, reviewer capacity
Reader	Primary, anonymous	Discover, read, cite and download published papers
CBAS college leadership	Secondary	Editorial throughput and workload analytics
Indexing services	External system	Harvest metadata via OAI-PMH

Non-human stakeholders: the AWS platform (cost and operability, section 13) and future maintainers, for whom the technical debt register (section 12) is written.

A single account may hold multiple roles simultaneously — a deliberate design choice (design specification section 3, SRS section 2.3) with a direct security consequence: the authorisation layer does not currently prevent an Author–Reviewer dual-role holder from reviewing their own manuscript (TD-02, section 12).

5. Requirements analysis

Requirements were elicited from the problem statement (section 2) rather than from structured interviews with SDJ's editorial office, and no real SDJ submissions or reviewer data were available to the project (SRS section 2.6) — a stated constraint, not a hidden one. The analysis proceeded in four steps, each of which left an artefact this consolidation points to rather than repeats:

- 1. Actor and use-case inventory.** Eight stakeholder classes (section 4) were mapped to 25 use cases spanning identity, submission, screening, review, decision, publication, archive and administration (effort estimation section 3; SRS section 4.2).
- 2. Functional requirements, stated testably.** Each of 28 functional requirements (plus one, FR-25a, added during this analysis — see below) is written as "the system shall ...", with explicit preconditions, postconditions and an acceptance criterion, following IEEE 830-1998 and ISO/IEC/IEEE 29148:2018 (SRS section 1.1, section 3.1).
- 3. Non-functional requirements, each with a verification method.** Seventeen NFRs across security, integrity, performance, reliability, availability, usability, maintainability, observability, portability and compliance, each naming the test or check that verifies it (SRS section 3.2) — "non-functional" is deliberately not allowed to mean "unverifiable."
- 4. Prioritisation against an effort estimate, not against intuition.** MoSCoW priorities were assigned per requirement in the design specification (section 5) and made authoritative once the effort estimate (section 7 below) confirmed the full system was roughly two orders of magnitude larger than the 48-hour window — the cut was a consequence of measurement, not a guess made in parallel with it.

A requirement discovered during analysis, not invented for the report. Comparing the implemented lifecycle guard (`transitions.py` , which permits `WITHDRAWN` as a target state) against the implemented authorisation layer (`policies.py` , which has no action gating who may invoke it) surfaced a live authorisation

gap the original FR-25 did not anticipate. This was recorded as **FR-25a**, marked NEW, rather than silently patched — a demonstration that requirements analysis did not stop once the SRS was signed off (SRS section 3.1, Group E).

The full requirement set, with preconditions, postconditions and acceptance criteria, is in `docs/02-srs.md` section 3; this section summarises the *process*, that document contains the *content*.

6. Software Requirements Specification (SRS)

The authoritative SRS is `docs/02-srs.md`, conformant to IEEE 830-1998 and ISO/IEC/IEEE 29148:2018. It is structured in seven sections: introduction and scope (section 1), overall description and operating environment (section 2), specific functional and non-functional requirements (section 3), system models — lifecycle, actor mapping, authorisation matrix (section 4), the requirements traceability matrix (section 5), MoSCoW prioritisation (section 6), and constraints/limitations (section 7). This section summarises its content and states plainly where it is now stale relative to the delivered system (section 10.4).

6.1 Requirement counts

Category	Count	Priority split
Functional requirements	34 (FR-01...FR-34) + FR-25a	19 Must, 12 Should, 3 Could
Non-functional requirements	17 (NFR-01...NFR-17)	Security (6), Integrity (1), Performance (2), Reliability (1), Availability (1), Usability (1), Maintainability (2), Observability (1), Portability (1), Compliance (1)

FR-29 to FR-34 cover capabilities the first requirement set never anticipated: article processing charges, the administrator console, decision certificates, anonymisation preflight, review deadlines, and self-service author registration. They were written after their implementation, by reading the delivered routes and tests, so they describe what exists rather than what was intended.

6.2 The manuscript lifecycle is the SRS's executable core

SRS section 4.1 transcribes the lifecycle directly from `backend/src/ugjcs/domain/transitions.py`'s `LEGAL_TRANSITIONS` mapping — stated explicitly as the tested source of truth, not the narrative diagram in the design specification. The full state table and diagram are reproduced and explained in section 8.2 of this document, since the lifecycle is equally central to system analysis.

6.3 A recorded, load-bearing disagreement between two authoritative documents

SRS section 4.1.3 records that the design specification's own lifecycle diagram (section 6.2 of that document) draws `RESUBMITTED` flowing directly into `REVIEWS_COMPLETE`, implying a resubmission automatically closes the review round. The implemented `LEGAL_TRANSITIONS` table instead routes `RESUBMITTED` to **either** `UNDER_REVIEW` **or** `UNDER_SCREENING`, at editorial discretion, and never closes the round by itself. Per this project's stated rule — the implementation governs, and the disagreement is recorded rather than hidden — the SRS and this document both follow the code.

6.4 Traceability matrix — read honestly, not optimistically

SRS section 5 has been reconciled against the running code twice: once when the `api/` layer and frontend first existed, and again after the two feature waves. The second pass corrected rows that had drifted in the opposite direction, recording expertise-ranked matching, structured reviews, public document download, audit-chain exposure, citation export and analytics as unbuilt when all six had shipped. A matrix that understates the system is as much a defect as one that overstates it.

Current totals: 24 of 35 requirement lines implemented and tested end to end, 5 partially implemented with a named remainder, 6 unbuilt. `docs/05-api-contract.md` and `docs/06-testing-report.md` are the authorities on what is reachable by route and proven by test; section 10.4 reconciles the three.

6.5 Authorisation matrix — the register's origin point

SRS section 4.3 derives the authorisation matrix directly from `policies.py`'s `_ROLE_GRANTS` and `_OWNERSHIP_ACTIONS`, and names the two critical gaps that became TD-02 and TD-03 (section 12) before the technical debt register existed as a separate document — the SRS is where those gaps were first written down.

7. Software effort estimation

Full arithmetic, every intermediate figure, and the estimation method are in `docs/03-effort-estimation.md`. This section summarises the method, the headline figures and why they matter, without repeating the derivation.

7.1 Method

Use Case Points (UCP) is the primary technique — actor and use-case weights are read directly off the functional requirements table rather than guessed. **COCOMO II Early Design** serves as an independent cross-check on different inputs (source lines of code, process/product ratings) rather than actor/transaction counts. Agreement between two methods driven by different inputs is stronger evidence than precision from either alone.

7.2 Headline figures

Quantity	Value	Basis
Unadjusted actor weight (UAW)	19	6 GUI actors × 3 + 1 API actor × 1
Unadjusted use-case weight (UUCW)	225	25 use cases, Karner transaction-count classes
UUCP	244	UAW + UUCW
Technical complexity factor (TCF)	1.105	13 technical factors, $\Sigma T = 50.5$
Environmental complexity factor (ECF)	0.605	8 environmental factors, $\Sigma E = 26.5$
UCP (full system)	163.1	UUCP × TCF × ECF
Effort, full system	≈ 3,262 person-hours (≈ 1.8 person-years)	UCP × PF (20 h/UCP)
UCP (Must-have subset, UC1-UC18)	125.7	Same method, 18 use cases only
Effort, Must-have subset	≈ 2,514 person-hours	UCP × 20
COCOMO II Early Design (full system)	≈ 7,170 person-hours (47.2 person-months)	12 KSLOC, $\Sigma SF = 14.36$, $\prod EM = 1.1699$

PF = 20 h/UCP follows Karner's productivity-factor rule: the count of E1–E6 rated below 3, plus the count of E7–E8 rated above 3, sums to 0 here — within the ≤ 2 threshold for PF = 20, not a default assumption (effort estimation section 6).

7.3 Reconciliation, not convergence

The two methods differ by roughly $2.2\times$ ($7,170 \div 3,262$). This is explained, not dismissed: COCOMO II is calibrated on projects carrying formal verification and management overhead this project does not incur, and its SCED (schedule-compression) penalty compounds multiplicatively with RCPX and PDIF for an extreme-compression solo project. UCP, conversely, counts only externally visible actor transactions and is structurally blind to platform work — Terraform, CI/CD, the hash chain's internals — of which this project has a disproportionate amount relative to its use-case count. The two methods **bound the answer from the same side**: both place the full system in the one-to-four-person-year range, nearly two orders of magnitude beyond the 48-hour window. That agreement on *scale*, not either figure's precision, is what forces the MoSCoW cut.

7.4 The MoSCoW cut this estimate produced

Priority	Use cases	Decision
Must	UC1–UC18	Implemented to production quality
Should	UC19–UC23	Implemented only if Must-have work completed early
Could	UC24, UC25	Deferred; entered in the technical debt register

7.5 Why the realised effort does not invalidate the estimate

The Must-have estimate of 2,514 person-hours assumes Karner's PF = 20 h/UCP, calibrated on manual development. This build used AI-assisted development — Claude Code (Anthropic) pair-programming the implementation under the author's direction and review — which the estimation document treats as a change of development **method**, not merely of pace (effort estimation section 9). Three consequences follow, and each is load-bearing for how this document's conclusion (section 18) should be read:

- 1. UCP still correctly sized the problem.** What changes under AI assistance is the rate at which a chosen method converts problem-size into elapsed hours (the productivity factor), not the size of the problem itself. The MoSCoW cut in section 7.4 was decided before the realised productivity was known, and remains the correct basis for it.
- 2. The realised productivity factor is a local calibration, not a general claim** — evidence about this developer, this tool, this domain and this window, computed formally by the method in effort estimation section 10.3 (actual session-hours from commit history ÷ delivered UCP). It does not generalise to other developers or tools.
- 3. The gap between estimated and realised hours is not free capacity.** It is capacity that was not spent on activities the classical estimate priced in: test depth beyond the 85% coverage floor, architecture decision records and onboarding documentation, and security hardening beyond the NFR-01–NFR-06 baseline. This connects directly to the limitations recorded in section 17.

8. System analysis

This section and section 9 are original to this consolidation. Section 8 analyses the problem domain as modelled — the aggregates, the manuscript lifecycle as a state machine, the double-blind projection, the audit mechanism's design rationale, and the persisted data model. Section 9 covers the architectural design decisions that realise that analysis — the hexagonal layering, its two mechanically enforced contracts, the API design, and the deployment topology.

8.1 Domain model

`Manuscript` is the aggregate root; it owns its authorship, status, version and review count, and nothing outside the aggregate mutates them directly. `User` (implemented as `Account` in code — `domain/account.py`), `Issue` and the editorial event log are separate aggregates referenced by identity. Value objects (`domain/ids.py`) give every identifier — `TrackingCode`, `ManuscriptId`, `UserId` — its own type rather than a bare `UUID` or `str`, so a tracking code and a user id cannot be interchanged by a type error that would otherwise compile.

The eight domain modules and their responsibilities:


Module	Responsibility
<code>manuscript.py</code>	The <code>Manuscript</code> aggregate: lifecycle transitions, review-quorum counting, resubmission versioning
<code>transitions.py</code>	<code>LEGAL_TRANSITIONS</code> — the exhaustive state table and <code>assert_legal</code> guard

Module	Responsibility
policies.py	can(actor, action, resource) — role-based and ownership-based authorisation, deny by default
blinding.py	blind() — the structural double-blind projection (section 8.4)
hashchain.py	append / verify — the tamper-evident audit chain (section 8.5)
events.py	EditorialEvent — the audit record's canonical, hashable representation
account.py	Account aggregate: identity, roles, credentials (role vocabulary only — no registration flow in the domain layer)
enums.py, ids.py, errors.py	Shared vocabulary, typed identifiers, the DomainError hierarchy

8.2 The manuscript lifecycle as a state machine

The lifecycle below is transcribed from transitions.py's LEGAL_TRANSITIONS mapping — the executable, tested source of truth (SRS section 4.1), reproduced here because it is the single most important artefact in the system's analysis.

Source state	Legal targets	Terminal?
DRAFT	SUBMITTED	No
SUBMITTED	UNDER_SCREENING, WITHDRAWN	No
UNDER_SCREENING	DESK_REJECTED, UNDER_REVIEW, REVISION_REQUESTED, WITHDRAWN	No
UNDER_REVIEW	REVIEWS_COMPLETE, WITHDRAWN	No
REVIEWS_COMPLETE	ACCEPTED, REJECTED, REVISION_REQUESTED, WITHDRAWN	No
REVISION_REQUESTED	RESUBMITTED, WITHDRAWN	No
RESUBMITTED	UNDER_REVIEW, UNDER_SCREENING	No
ACCEPTED	SCHEDULED	No
SCHEDULED	PUBLISHED	No
DESK_REJECTED, REJECTED, PUBLISHED, WITHDRAWN	(none)	Yes



No transition is legal out of a terminal state — `assert_legal` raises `IllegalTransitionError` for any pair absent from `LEGAL_TRANSITIONS`, checked exhaustively over all 121 (source × target) status pairs by a property test (`test_lifecycle_permits_expected_transitions`, `test_lifecycle_forbids_shortcut_transitions`, testing report section 3.1–section 3.2). `ACCEPTED` and `SCHEDULED` are deliberately excluded from the withdrawable set: retracting an accepted paper is editorial retraction, with its own notice obligations, and is not modelled as author withdrawal.

8.3 Authorisation as an explicit relation, not scattered checks

`policies.can(actor, action, resource)` is a single function computing every access decision in the system, called as a FastAPI dependency so no route can omit it (section 9.3). It composes two kinds of rule: `_ROLE_GRANTS` (an action is available to any actor holding a given role — e.g. `SCREEN` requires `editor` or `editor_in_chief`) and `_OWNERSHIP_ACTIONS` (an action additionally requires the actor's identity to match a predicate over the resource — e.g. `RESUBMIT` and `WITHDRAW` require the actor to be the manuscript's corresponding author). Two of this matrix's known gaps — `REVIEW` granted on role alone with no ownership predicate excluding an author reviewing their own work — are TD-02 and TD-03 (section 12); they are gaps in the *matrix*, found by asking whether every action that should be ownership-checked is, not gaps in the mechanism that enforces it.

8.4 Double-blind integrity as a projection, not a filter

The core design decision: blinding is **structural**, not procedural. `blind(manuscript)` (`blinding.py`) does not filter a full `Manuscript` object and hope every caller remembers to call it — it returns a distinct type, `BlindedManuscript`, that has no author-identifying field at all:


```
@dataclass(frozen=True, slots=True)
class BlindedManuscript:
    tracking_code: str
    title: str
    abstract: str
    keywords: tuple[str, ...]
    version: int
    status: str
```

There is no field a future change could accidentally repopulate with an author id, because the type has nowhere to put one. This is verified by a property test asserting the author id is absent from the serialised projection for every generated title/abstract/keyword combination (testing report section 3.2), and by sentinel-based leak tests over every reviewer-facing endpoint (testing report section 3.5). The documented limit — `title`, `abstract` and `keywords` are copied **verbatim**, so self-identifying text in the body is not scrubbed — is TD-05 and is restated in Section 17.

8.5 The hash-chained audit log


Each `EditorialEvent` records its sequence number, type, payload, actor, timestamp, its predecessor's hash, and its own SHA-256 digest over `previous_hash || canonical_bytes(event)`:

```
def chain_hash(event: EditorialEvent, previous_hash: str) -> str:
    digest = hashlib.sha256()
    digest.update(previous_hash.encode("ascii"))
    digest.update(event.canonical_bytes())
    return digest.hexdigest()
```

`append` links a new event onto a chain tail in $O(1)$ — it reads only the last link, not the whole history (TD-08 records that this capability is unverified by a dedicated test). `verify` walks the chain from the genesis hash (`"0" × 64`) and raises `ChainBrokenError` at the first inconsistency in sequence, predecessor hash, or digest. This detects **alteration, reordering and removal within the chain** — it is deliberately, and explicitly in the module's own docstring, **not** a defence against tail truncation, a forged event appended through the legitimate API, or a wholly rebuilt history from genesis: all three require an external anchor the domain layer does not provide (TD-04, restated in section 17). A PostgreSQL trigger closes the most direct way to defeat this at the database layer for `UPDATE / DELETE`; a second, statement-level trigger closes `TRUNCATE`, which the first trigger did not — see section 11.4's account of how that gap was found.

8.6 Data model — entity-relationship diagram

The tables below are the delivered PostgreSQL schema (`backend/src/ugjcs/infrastructure/db/models.py`), not the fuller schema sketched in the design specification (section 8 of that document) — `reviews`, `editorial_decisions`, `issues`, `issue_papers`, `similarity_reports`, `notifications` and a dedicated roles table are specified there but not present in the delivered database; roles are a `user_roles` join table, and the review outcome is folded into `review_assignments`.



A design note worth stating rather than glossing over: `manuscript_authors.author_id` and `review_assignments.reviewer_id` are plain `UUID` columns, not declared foreign keys to `users.id` — referential integrity between authorship/reviewer assignment and the user table is an application-level guarantee, not a database-enforced one. `editorial_events.manuscript_id` is a declared foreign key with `ON DELETE RESTRICT`, so a manuscript with audit events cannot be deleted (verified: testing report section 3.3). `editorial_events` carries a unique constraint on `(manuscript_id, event_hash)` and no application code path issues `UPDATE` or `DELETE` against it; the database-level enforcement of that is described in section 8.5 and section 11.4.

9. System design


9.1 Architecture — hexagonal, backend

```

backend/src/ugjcs/
├── domain/      entities, value objects, events, state machine, policies
│               NO framework imports (enforced by import-linter, section 9.2)
├── application/ use-case services, port protocols, DTOs, unit of work
├── infrastructure/ SQLAlchemy repositories, S3 storage, JWT, Argon2, email
└── api/         FastAPI routers, request/response schemas, dependencies, wiring

```

Dependencies point inwards only. Concrete adapters (SQLAlchemy repositories, the S3 document store, JWT token service) implement port *protocols* declared in `application/ports.py` and are bound to them at a composition root (`api/wiring.py`), so the application and domain layers never import a concrete adapter. The payoff is directly demonstrable, not asserted: the entire domain test suite runs with no database, no network and no framework loaded (testing report section 1).



9.2 The two import-linter contracts, and why they are mechanical rather than aspirational

`.importlinter` at the backend root defines two contracts, both run as a CI gate in `make check` (testing report section 2):

- **layers** — `ugjcs.api` → `ugjcs.infrastructure` → `ugjcs.application` → `ugjcs.domain`. A layer may import only the layers named after it in this list; an import running the other way fails the build. This is what makes "dependencies point inwards" a property CI checks on every commit, not a convention a reviewer might miss.
- **domain-purity** — forbids `ugjcs.domain` from importing `fastapi`, `sqlalchemy`, `pydantic`, `boto3`, `arq`, `redis`, `httpx`, `requests`, and thirteen standard-library I/O/framework modules including `os`, `io`, `socket`, `logging`, `asyncio` and `subprocess`. A domain module that reaches for a timestamp via

`datetime.now()` is permitted; one that reaches for a file, a socket, or a framework decorator is not, and the build fails before a test even runs.


Both contracts are a direct, testable realisation of NFR-13 ("the domain layer imports no framework code") — the requirement is not merely documented, it is enforced the same way a type error is: at build time, unconditionally.

9.3 API design

The HTTP boundary is documented exhaustively in `docs/05-api-contract.md`; this section states its governing design decisions rather than repeating the endpoint table.

- **Bearer-token backend, cookie-sessioned frontend.** The FastAPI backend is a pure bearer-token API — short-lived JWT access tokens (HS256) plus rotating, hash-stored refresh tokens. The Next.js frontend is a Backend-For-Frontend: Route Handlers under `frontend/src/app/api/**` unseal an `httpOnly`, `Secure`, `SameSite=Lax` cookie and attach the bearer token server-side. No browser-side JavaScript ever holds a token. This exists specifically to avoid third-party cookie blocking between the `*.vercel.app` and `*.awsapprunner.com` origins (design specification section 7.2).
- **Errors as RFC 9457 Problem Details**, uniformly, from every endpoint — `{type, title, status, detail?, instance?}` served as `application/problem+json`, with a fixed exception→status mapping (`IllegalTransitionError`→409 , `AuthorizationDeniedError`→403 , etc.) so the frontend branches on structure, not on parsing a message string.
- **snake_case on the wire, unconditionally.** No camelCase translation layer exists anywhere; `ugjcs.domain.enums` values are copied verbatim into the frontend's TypeScript types. This was a considered choice, recorded as such in the API contract, not an oversight.
- **No pagination anywhere.** A deliberate scope decision for a demonstration-scale corpus, recorded so that adding it later is understood as an additive change, not a redesign.
- **Authorisation as a dependency, not a decorator convention.** Every route's `Action` is resolved through `policies.can()` as a FastAPI dependency, and `test_route_audit.py` walks the *live* route table asserting every non-public route carries one — a defect a hand-maintained checklist would eventually miss (testing report section 1).

9.4 Sequence — submit to publish



This flow is exercised end to end by the acceptance script in testing report section 5, and the `schedule / publish` steps are the specific ones that were, for a period, implemented and unit-tested but reachable by no route — the finding retold in section 11.4 and cross-referenced from Section 10.4.

9.5 Deployment topology

Covered fully in section 13; the topology is a component of the design and is diagrammed there once alongside its discrepancy from the design specification (TD-14) rather than twice.

10. Implementation

10.1 Technology stack

Layer	Choice	Rationale (design specification section 7.4)
Backend	FastAPI, Python 3.13	Native async I/O; Pydantic v2 gives validation and OpenAPI from one type
ORM	SQLAlchemy 2.0 + Alembic	Mature migrations; mapping style keeps domain classes framework-free
Database	PostgreSQL 16	Native full-text search and JSONB avoid a separate search engine
Storage	S3	Durable, private, pre-signed access only

Layer	Choice	Rationale (design specification section 7.4)
Frontend	Next.js 15, TypeScript, Tailwind	Static public pages plus a server-side BFF in one deployment
Testing	pytest, Hypothesis, testcontainers, Vitest	Layer-appropriate verification (section 11)
IaC	Terraform	Reproducible, reviewable, destroyable

10.2 What is implemented, by layer

Domain (`backend/src/ugjcs/domain/` , 812 lines across 11 modules) — complete: the manuscript lifecycle, RBAC policy, the double-blind projection, the hash chain, the account aggregate, typed identifiers. This is the layer with no external dependency and the highest test confidence (section 11).

Application (`backend/src/ugjcs/application/` , 487 lines) — port protocols (`ports.py` , 236 lines), identity-related use-case orchestration (`identity.py` , 186 lines) and document handling (`documents.py` , 65 lines).

Infrastructure (`backend/src/ugjcs/infrastructure/`) — SQLAlchemy repositories and unit of work (`infrastructure/db/`), an S3-backed document store with a demo-PDF generator for seeding (`infrastructure/storage/`), Argon2 password hashing and JWT issuance (`infrastructure/security/`), and a logging-only email sender (`infrastructure/email/logging_sender.py` — no live transactional provider is wired in this build, a limitation restated in section 17).

API (`backend/src/ugjcs/api/` , 966 lines) — five routers (`auth` , `manuscripts` , `editorial` , `reviews` , `archive`), RFC 9457 error mapping (`errors.py`), and the composition-root wiring (`wiring.py`) that binds concrete adapters to the application's port protocols. This layer, and the frontend below it, are what SRS section 5 (written earlier in the project) records as "Planned" — Section 10.4 reconciles that.

Frontend (`frontend/src/app/`) — a public route group (`(public)/search` , `(public)/papers`) requiring no authentication, and role-scoped authenticated routes (`author/` , `author/submit/` , `author/[trackingCode]/` , `editor/` , `editor/[trackingCode]/` , `reviewer/` , `reviewer/[trackingCode]/` , `login/`), plus a `frontend/src/app/api/**` BFF layer mirroring the backend's resource groups. There is no dedicated Administrator UI — role management is verified only at the policy-test level (testing report section 5, section 6), consistent with FR-03's grant existing in `policies.py` without a corresponding administration screen.

10.3 Domain code, as evidence rather than assertion

The lifecycle guard and the hash chain shown in section 8.2 and section 8.5 are reproduced directly from source, not paraphrased, specifically so a reader can check this document's claims against the file it names. The same discipline extends to the authorisation layer (section 8.3) and the blinded projection (section 8.4): every mechanism described in this section has a named module, and every claim about its behaviour is backed by a named test in `docs/06-testing-report.md`.

10.4 Reconciling three documents that were written at different points in the build

This consolidation surfaced a genuine discrepancy between three of its own source documents, stated here rather than resolved silently in one direction, consistent with section 6.3's handling of the SRS-versus-specification lifecycle disagreement:

- **docs/02-srs.md section 5** (earliest) records almost the entire functional requirement set as "Planned," on the stated grounds that no `api/` directory and no frontend existed in the repository at the time it was written.
- **docs/05-api-contract.md** (written once the API and frontend existed) documents a working `/auth` , `/manuscripts` , `/editorial` , `/reviews` and `/archive` surface — but explicitly states that `Action.PUBLISH / Manuscript.schedule / Manuscript.publish` have **no corresponding route**: "publication into the archive happens outside the HTTP boundary this plan builds."

- `docs/06-testing-report.md` **section 3.4 and section 4.5** (latest) records passing system tests named `test_the_editor_in_chief_can_schedule_an_accepted_manuscript` and `test_the_editor_in_chief_can_publish_a_scheduled_manuscript`, and states directly that these two routes, along with manuscript resubmission, were **added after** being found reachable by no route during manual use of the deployed system.

Read together rather than in isolation, these three documents describe the same system at three successive points in its build, not three inconsistent descriptions of one static state: the SRS predates the API layer; the API contract predates the schedule/publish routes; the testing report postdates their addition and is the most current account of what the deployed system can do. This document treats the testing report as authoritative on current reachability, the API contract as authoritative on wire format for the routes it does describe, and the SRS as authoritative on requirement *content and traceability method* rather than current build status. A reader integrating all three should not conclude either that publication is unreachable (05, superseded) or that the majority of FRs remain unbuilt (02 section 5, superseded) — section 11.4's account of finding and closing that exact gap is the more current and more informative story.

11. Testing

Full test-case tables, defect narratives and the UAT script are in `docs/06-testing-report.md`. This section summarises the strategy, headline numbers, and the finding this project treats as its most important methodological result (section 11.7, elaborated for the whole document in the introduction's through-line).

11.1 Strategy — layered to match the architecture

Layer	What it tests	Tooling
Unit — domain	Pure logic, no database/HTTP/mocks	pytest
Property-based	Invariants over generated inputs (lifecycle, hash chain, blinding)	Hypothesis
Unit — application/API/db/security	Adapter logic in isolation	pytest, fakes
Integration	Real PostgreSQL via <code>testcontainers</code> — trigger firing, <code>timestampz</code> normalisation, FK behaviour	pytest, testcontainers
Contract	Architectural rules, not behaviour	import-linter, <code>mypy --strict</code>
Route audit	Every non-public route carries an authorisation dependency	pytest walking the live route table
End-to-end / manual acceptance	Scripted, role-scoped scenarios against the live deployment	Playwright-driven browser session

11.2 Headline numbers (2026-08-14, CI run 31791711326)

```
make check      → 402 passed, 84 deselected. Coverage: 90.03% (gate: 85%)
make integration → 84 passed, 402 deselected
npx vitest run  → 79 passed (20 files)
```

486 backend tests and 79 frontend tests, 565 in total. The backend figures are taken from a CI run rather than a local one, because the 84 integration tests stand up a real PostgreSQL 16 through `testcontainers` and cannot run on a machine without Docker; there they error at setup instead of failing, which is a more honest signal than a false pass but not a result worth citing.

Coverage is gated on `domain` and `application` only. That is a documented decision rather than an oversight: measuring infrastructure coverage in a run that has excluded the infrastructure tests would credit incidental import-time coverage, not genuine exercise.

11.3 CI gates

`.github/workflows/backend-ci.yml` runs two jobs on every push/PR to `main / master`: **check** (ruff lint, ruff format, mypy strict, import-linter, unit suite at the 85% gate) and **integration** (a real `postgres:16` service container, Alembic migration applied **up, then down, then up again** to verify reversibility, then the integration suite). Neither job passing is optional for merge.

11.4 Defects found — the section with the most to learn from

Six defects are recorded in testing report section 4, none found by a coverage number or a green test run alone:

#	Defect	Found by
1	Deleting the hash chain's chaining line left all tests passing at 100% coverage	Manual mutation testing
2	<code>canonical_bytes()</code> hashed a UTC-offset-bearing timestamp; PostgreSQL normalises offsets on storage, so a round trip would false-positive <code>verify()</code> as tampered	Reasoning about the storage boundary, not any test
3	A row-level trigger blocked <code>UPDATE / DELETE</code> on the audit log but PostgreSQL never fires row-level triggers on <code>TRUNCATE</code> — one statement erased the log with no error	Asking a broader question of an already-verified control
4	<code>coverage.py</code> does not model a ternary as a two-arm branch — "0 branches missing" while one arm never ran	Reading logic against the coverage report, not trusting it
5	Three domain lifecycle methods (<code>resubmit</code> , <code>schedule</code> , <code>publish</code>) were fully implemented and unit-tested but reachable by no API route	The owner using the deployed system as an actual actor would
6	A container held IAM permission to reach S3 but no network route to it — uploads hung until the health check killed the instance	Exercising the deployed upload feature against real infrastructure

Defects 1–4 were found by review of code that had already passed linting, strict typing, an architecture contract and a 100%-covered test suite; defects 5–6 were found only by using the running system. This is the single clearest piece of evidence this project produced about software engineering practice, and it is why the technical debt register (section 12) and this report both state it directly rather than let it sit implicit in a defect table.

11.5 User acceptance testing

Run as scripted, role-scoped scenarios against the live deployment (`https://ugjcs-frontend.vercel.app`, backed by `https://tsxsbf9rzp.us-east-1.awsapprunner.com`), with five named judge accounts (author, reviewer, editor, editor-in-chief, administrator — see Section 14). Three scenarios were exercised live for the testing report (author login and dashboard, unauthenticated archive access, cross-role access denial redirecting to `/` rather than `/login` — a deliberate distinction in `frontend/middleware.ts`); the remaining role-scoped scenarios are documented against their automated backstop test.

11.6 What testing did not cover

Stated plainly in testing report section 6, not omitted: no load or performance testing (NFR-08/NFR-09 are unverified against the running system); no automated security scanning in CI (no SAST, DAST or dependency vulnerability scan); no mutation testing in CI (section 11.4's finding #1 was a one-off manual pass); no browser-matrix testing (one Chromium-based session only); no committed automated end-to-end suite (`@playwright/test` is a declared dependency with no spec files committed). Each is restated as a limitation in section 17.

11.7 Evaluation

The testing report's own conclusion, reached independently from the test evidence, matches the technical debt register's closing observation exactly: **automated gates establish a floor and catch regressions; they did not find the defects that mattered most**. Every serious defect was found by a human or an agent

reading code against what it claimed to do, by mutation testing designed specifically to distrust the coverage figure, or by using the running system as an actual actor would. This is elaborated once, fully, in section 12.6, since the two documents converge on one finding rather than two.

12. Technical debt

The full register — fourteen entries, each with Debt → Cause → Impact → Priority → Proposed resolution — is `docs/04-technical-debt-register.md`. This section summarises it by priority and draws out its methodological finding.

12.1 Summary

Priority	Count	Entries
Critical	3	TD-01, TD-02, TD-03
Scheduled	6	TD-04, TD-05, TD-06, TD-07, TD-08, TD-14
Acceptable	3	TD-09, TD-10, TD-11
Resolved, retained as a record	2	TD-12, TD-13

12.2 Critical — must be resolved before real users

ID	Debt	Impact
TD-01	AWS access uses root account credentials	Cannot be scoped, rotated, or revoked without disrupting the whole account; compromise is unrecoverable within it. Mitigated: root keys are used only from the developer's workstation, never stored as a CI secret
TD-02	<code>Action.REVIEW</code> is granted on the <code>REVIEWER</code> role alone, with no per-manuscript predicate	An Author–Reviewer dual-role holder is not prevented from reviewing their own work — the central conflict-of-interest failure for a double-blind journal
TD-03	Submitted reviews are counted, not identity-checked against an accepted assignment	One reviewer calling submit twice reaches quorum alone and can close a review round unilaterally

12.3 Scheduled — accepted now, with a named repayment point

ID	Debt	Repayment point
TD-04	The audit chain has no external anchor — tail truncation is undetectable by the application alone	Next release after deployment
TD-05	Blinding does not scrub the manuscript body — <code>title / abstract / keywords</code> are verbatim	Screening surfaces name matches to the editor; automated redaction is future evolution, not near-term
TD-06	The editorial event log has no blinded projection (<code>actor_id</code> , rationale text carried in full)	Before any reviewer-facing audit view is built
TD-07	No <code>Action</code> connects <code>blind()</code> to <code>policies.can()</code> — an adapter must remember to call it	Alongside TD-02
TD-08	The tail-append capability (section 8.5) is unverified by a dedicated test	With the persistence work that first exploits it
TD-14	Deployed infrastructure is App Runner, not the ECS/ALB/CloudFront topology specified	After submission — repayable once the pre-submission document/feature backlog is clear (section 13.3)

12.4 Acceptable — a conscious, revisitable trade-off

ID	Debt	Condition that would change the judgement
TD-09	A hybrid event log (materialised status + append-only log) rather than full event sourcing	If projections multiply beyond one, or replay-to-a-past-state becomes a requirement
TD-10	The coverage gate sits at 85% while the code delivers 88–100%	Cheap to raise; deliberately not set to the exact current figure so a legitimate refactor doesn't fail the build
TD-11	Coverage is a weak signal — the register's own evidence (section 12.6)	Compensated by mutation testing and review, not by the gate alone

12.5 Resolved, retained as a record

TD-12 (a UTC-offset timestamp representation would have produced a false tamper alert) and TD-13 (TRUNCATE bypassed the row-level append-only trigger) are both closed, but kept in the register rather than deleted, because the *class* of defect — and how each was found — is the point (section 12.6, Section 11.4).

12.6 The register's own conclusion, and this document's central finding

`docs/04-technical-debt-register.md` closes with a statement worth repeating verbatim in spirit rather than only citing: **ten of its fourteen entries, and every serious defect in `docs/06-testing-report.md` section 4, were found by independent review of code that had already passed every automated gate available in this project — linting, strict type checking, an architecture contract, and a full test suite at 100% coverage — or by a person using the running system.** Mutation testing showed the hash chain was, for a period, unprotected by any test that would notice its defining property being deleted. TRUNCATE bypassed a trigger that had been "confirmed firing against a live database" against the two statement types someone thought to check. Three lifecycle methods were implemented, unit-tested, type-checked and covered at 100%, and were still dead code from the system's perspective because no route called them. Passing every gate this project has was, in every one of these cases, necessary and not sufficient. That is this project's most defensible finding about software engineering practice, and it is the reason section 15's maintenance strategy treats the debt register as a *live* input to a repayment schedule rather than an inventory closed at submission.

12.7 Repayment sequence

TD-01 before any further infrastructure is provisioned. TD-02, TD-03 and TD-07 are one piece of work — all three are consequences of reviewer assignment not existing as a first-class entity — and should be repaid together in the release that introduces it. TD-04 follows deployment. TD-05, TD-06 and TD-08 are independent and may be scheduled by convenience. TD-14 is repayable only after submission and is not on the pre-viva critical path.

13. Deployment


13.1 What is actually running

Component	Where	Notes
Frontend	Vercel, <code>ugjcs-frontend.vercel.app</code>	Next.js 15, Git-integrated deploys
Backend API	AWS App Runner, <code>tsxsb9rzp.us-east-1.awsapprunner.com</code>	256 CPU units / 512 MB, VPC connector for egress
Database	RDS PostgreSQL 16, <code>db.t4g.micro</code>	Private, <code>publicly_accessible = false</code> , default VPC subnets, single-AZ
Object storage	S3, <code>ugjcs-manuscripts-<random></code>	All public access blocked; versioned; SSE-AES256; documents reached only via pre-signed URLs

Component	Where	Notes
Secrets	AWS Secrets Manager	ugjcs/jwt-secret , ugjcs/database-url — both random_password-generated, never a literal in Terraform
Container registry	ECR	ugjcs-backend repository

Infrastructure is defined entirely in `infra/` (Terraform): `network.tf`, `rds.tf`, `s3.tf`, `secrets.tf`, `iam.tf`, `apprunner.tf` / `apprunner_service.tf`, `ecr.tf`, `security_groups.tf`, `s3_endpoint.tf`, `outputs.tf`, `providers.tf`, `variables.tf` — satisfying NFR-16 ("the entire infrastructure is reproducible from code").

13.2 Deployment topology



App Runner supplies its own `*.awsapprunner.com` TLS endpoint, which is what makes an HTTPS frontend able to reach the backend without a registered domain — the same guarantee the design specification's CloudFront strategy (section 7.3 of that document) was chosen for, over the same container image.

13.3 The gap between this topology and the specified one — TD-14

The design specification (section 7.3) specifies `Reader → Vercel → CloudFront → ALB → ECS Fargate (api) → RDS / S3 / Redis → ECS Fargate (worker)`. What is deployed is `Reader → Vercel → App Runner → RDS / S3`, with no CloudFront, no ALB, no ECS, no Redis, and no worker service. This is recorded in the technical debt register as **TD-14**, deliberate and scheduled: provisioning the ECS/ALB/CloudFront stack — a target group, listener rules, a CloudFront distribution, task definitions, and the IAM wiring between all of it — was measured at 4–6 hours against a 48-hour budget that, at the point the trade-off was made, still owed a working API, a working frontend, and five accompanying documents. No functional capability is lost by the substitution: App Runner supplies the identical trusted-TLS-without-a-registered-domain guarantee CloudFront was chosen for, over the same image, with less to operate and less to tear down. The absence of Redis and a worker service means the asynchronous submission-processing pipeline described in the design specification (section 10.2 of that document — text extraction, similarity screening, anonymised derivative generation) is not deployed either; this is restated in section 16 and section 17 rather than left implicit in the infrastructure diagram alone.

13.4 CI/CD

`.github/workflows/backend-ci.yml` gates every push and pull request (section 11.3). Deployment itself is run **locally**, not from CI — a deliberate mitigation for TD-01: AWS root credentials are used only from the developer's workstation and are specifically not stored as a GitHub Actions secret, which meaningfully reduces exposure relative to a CI secret store, though it is a mitigation, not a resolution of the underlying debt. The frontend deploys through Vercel's Git integration on every push to the default branch.

13.5 Operating cost

Targeted at USD 35–55/month against AWS and Vercel free/low-cost tiers (design specification section 13, Section 16). App Runner at the smallest instance size (0.25 vCPU / 0.5 GB) is a deliberate cost decision recorded directly in `infra/apprunner_service.tf`'s own comments: roughly USD 14/month against ~USD 57 for the next tier up, since App Runner bills provisioned capacity continuously while the health check keeps the instance active.

14. User manual

`docs/07-user-manual.md` is being written concurrently with this document and is the authoritative source once complete; this section summarises the system's user-facing surface directly from the implemented routes and roles (section 10.2, `docs/05-api-contract.md`) so that a reader has a usable account of "how to use the system" even before that document lands. Where the two disagree once `07-user-manual.md` exists, that document governs.

14.1 Accounts and roles

Five roles exist (`Role` enum, `domain/enums.py`): `author`, `reviewer`, `editor`, `editor_in_chief`, `administrator`. A single account may hold several roles at once (section 4). Judge accounts used for acceptance testing (testing report section 5):

Role	Account
Author	<code>author@sdj.test</code>
Reviewer	<code>reviewer@sdj.test</code>
Editor	<code>editor@sdj.test</code>
Editor-in-Chief	<code>eic@sdj.test</code>
Administrator	<code>admin@sdj.test</code>

14.2 As an author

Log in at `/login`; the session lands on `/author`, listing your own submissions with status (never another author's). Submit a new manuscript at `/author/submit` with title, abstract, keywords and optional co-authors (JSON submission — there is no file-upload path in the delivered domain; see section 17). View a specific manuscript and its status at `/author/[trackingCode]`. If a manuscript is returned with `revision_requested`, resubmit from the same page; only the corresponding author may do so (`Action.RESUBMIT`, ownership-checked). A manuscript may be withdrawn from any of `submitted`, `under_screening`, `under_review`, `reviews_complete` or `revision_requested`, again corresponding-author-only.

14.3 As a reviewer

`GET /reviews/mine` (surfaced at `/reviewer/[trackingCode]`) lists only manuscripts you are assigned to review, in the **blinded** form: title, abstract, keywords, version and status — no author name, affiliation or identifier of any kind (section 8.4). Submit a recommendation and free-text comments; there are no per-criterion scores on the delivered wire format (a gap from the design specification's fuller `Review` model, noted in `docs/05-api-contract.md` section 8).

14.4 As an editor

`/editor` lists the screening queue (manuscripts in `submitted`). Screen a submission to move it into `under_screening`; from there, send to review, request pre-review changes, or desk-reject. Assign a reviewer directly by id — there is no candidate-recommendation UI in the delivered system (FR-08 is planned, not built; section 17). Record a decision once the review quorum is met.

14.5 As Editor-in-Chief

All Editor capability, plus scheduling an accepted manuscript into an issue and publishing it — both denied to a plain Editor (verified: `test_a_plain_editor_cannot_schedule`, `test_a_plain_editor_cannot_publish`, testing report section 3.4).

14.6 As an Administrator

Role management (`Action.MANAGE_USERS`) is granted and denied correctly at the policy layer, but has **no frontend surface** in this build — verified only by `test_administrator_may_manage_users` at the policy-test level (testing report section 5, section 6). There is no `/admin` route to walk through.

14.7 As a reader (no account required)

`/search` and `/(public)/papers` require no authentication. Browse and search published papers; download the original PDF of any published paper. Citation export (BibTeX/RIS) and OAI-PMH harvesting are specified (design specification section 10.3) but not present in the delivered API surface (`docs/05-api-contract.md` section 6 — no `/archive` citation-export or `/oai` route exists).

15. Maintenance strategy

No maintenance has yet occurred — the system is newly deployed. This section states the strategy that would govern maintenance from this point, organised by the four classical ISO/IEC 14764 maintenance categories, and states plainly what monitoring exists today and what does not.

15.1 Corrective maintenance

Defects are triaged by the same severity language the technical debt register already uses (Critical / Scheduled / Acceptable, section 12), so a newly discovered defect slots into the existing repayment sequence rather than requiring a new taxonomy. Given section 11.7's finding — that automated gates did not catch the defects that mattered — corrective maintenance should not rely on CI alone even after TD items are closed: a periodic manual review pass, and mutation testing once it exists in CI (currently absent, TD-11), are both load-bearing parts of this strategy, not optional extras. Every corrective fix should add the regression test that would have caught it, following the pattern already established for TD-12 and TD-13 (section 11.4).

15.2 Adaptive maintenance

Two adaptive pressures are already named and scheduled rather than hypothetical: **TD-14** (migrating from App Runner to the specified ECS/ALB/CloudFront topology, section 13.3) and the introduction of reviewer assignment as a first-class entity (the work that closes TD-02, TD-03 and TD-07 together, section 12.7). Both are adaptive in the ISO sense — changing the system to fit an environment (a scaled deployment target; a real reviewer-matching subsystem) rather than fixing a defect in the current one.

15.3 Perfective maintenance

Ordered by the repayment sequence in section 12.7: TD-01 (least-privilege IAM) first, as it blocks admitting real users; then the TD-02/TD-03/TD-07 reviewer-assignment work; then TD-04 (an external audit anchor); then TD-05/TD-06/TD-08 as convenient. Raising the coverage gate (TD-10) is low-cost and should be revisited whenever coverage is observed to drift, not on a fixed schedule.

15.4 Preventive maintenance

Dependency and security updates are currently manual and ad hoc: `pip-audit` and `npm audit` are named as intended CI gates in the design specification (section 11) but are **not** present in `.github/workflows/backend-ci.yml` today (testing report section 6) — this is itself a preventive-maintenance gap, not only a testing gap, and should be the first thing added to CI once the critical-priority items in section 12.2 are closed. Systematic mutation testing (`mutmut` or `cosmic-ray`, TD-11) belongs in the same category: a preventive control against the specific class of defect section 11.4 shows this project's existing gates cannot see.

15.5 What monitoring exists, and what does not

Exists: liveness and readiness probes (`/health` , `/ready` , NFR-11), consumed by App Runner's own health check to replace an unresponsive instance automatically — this is precisely the mechanism that surfaced defect section 11.4 #6 (a hung upload killed by the health check, not caught by any application-level alert).

Does not exist: no structured-log aggregation or trace export is deployed, despite NFR-15 specifying one (SRS section 3.2) — App Runner's default logging is what is actually available, not the correlation-ID-carrying structured JSON the requirement describes. No alerting is configured beyond the platform's own instance-replacement behaviour. No dashboard exists for the editorial analytics FR-24 specifies (that FR is itself unbuilt, section 10.4). No performance monitoring exists (NFR-08/NFR-09 are unverified in production, section 11.6). A maintainer inheriting this system should treat "nothing paged me" as meaning nothing detectable paged, not as evidence of correct operation — the same caution section 12.6 draws about automated gates applies to monitoring that has never been exercised against a real incident.

16. Future evolution

Drawn from the technical debt register's scheduled items and the design specification's own future-work section (section 17 of that document), organised by size rather than by document of origin:

Reviewer matching with expertise scoring. The design specification (section 10.1) specifies a TF-IDF vocabulary over reviewer expertise and manuscript text, hard exclusions (author, affiliation match, prior decline, unavailability, capacity), and a Hungarian-algorithm (`scipy.optimize.linear_sum_assignment`) global assignment, editor-overridable. None of this is built; reviewer assignment today is a persistence-only record with no matching, no invitation lifecycle, and no conflict-of-interest check (`docs/05-api-contract.md` section 8). This is also the piece of work that retires TD-02, TD-03 and TD-07 together (section 12.7) — the highest-leverage single addition against the current debt register.

The asynchronous processing pipeline. Text extraction, MinHash/LSH similarity screening against the internal corpus, and metadata-stripped anonymised-derivative generation, enqueued on upload and keyed by content checksum for idempotency (design specification section 10.2). Requires the Redis/ARQ worker component that TD-14 also notes is absent from the current deployment (section 13.3) — these two gaps compound, and closing the deployment gap first is a precondition for this one.

OAI-PMH and citation export. `Identify` / `ListMetadataFormats` / `ListIdentifiers` / `ListRecords` / `GetRecord` over Dublin Core with resumption tokens (FR-22), plus BibTeX/RIS export (FR-21) — both Should-have priority (section 7.4), specified, and absent from the delivered `/archive` surface.

Full event sourcing. TD-09 records the current hybrid (materialised status + append-only event log) as an acceptable, revisitable trade-off. The stated condition for revisiting it: if projections multiply beyond the current single materialised view, or replaying history to a past state becomes a requirement, full event sourcing with rebuildable projections becomes worth its cost.

Multi-journal tenancy. Explicitly out of scope for this build (design specification section 4.2), but the codebase separates journal-configuration data from platform logic as a deliberate seam (design specification Appendix A) specifically so this remains reachable without restructuring — future evolution here is closer to "activate a seam" than "redesign."

Smaller items, in the same register: an external, tamper-resistant anchor for the audit chain (TD-04); automated detection or redaction of self-identifying text in manuscript bodies beyond metadata stripping (TD-05, judged a research problem in its own right — a bad redaction that leaks one name would undermine the guarantee more than not attempting it); a `BlindedEvent` projection for a future reviewer-facing audit view (TD-06); embedding-based reviewer matching to replace TF-IDF, ORCID authentication, real Crossref DOI registration, a production typesetting/galley pipeline, reviewer reputation modelling, blue-green deployment, and read-replica scaling for the archive (design specification section 17, none yet begun).

17. Limitations

Stated plainly, each cross-referenced to where it is verified or recorded, following the same convention as SRS section 7 and the technical debt register.

- **This system was built in 48 hours by one developer with AI assistance.** Every limitation below is a direct consequence of that constraint, not of an unconsidered choice; section 7.5 explains why the resulting scope decisions were nonetheless estimation-driven rather than ad hoc.
 - **Anonymisation strips PDF metadata but not author names typed into the body.** `blind()` (section 8.4) is structurally guaranteed to omit author fields from the type a reviewer receives — but `title`, `abstract` and `keywords` are copied verbatim. An author who writes their own name into the title, or an abstract that reads "extending our earlier work in [Obeng 2025]," reaches the reviewer unchanged. Double-blind integrity therefore depends partly on author compliance with submission guidance, not entirely on what the system enforces (TD-05).
 - **The audit chain has no external anchor.** Hash chaining detects alteration, reordering and removal *within* the chain (section 8.5), and a database trigger blocks direct tampering — but truncation of the chain's tail, or a forged event appended through the legitimate API, is **undetectable by the application alone**, because there is no periodically published, independently held checkpoint of the latest `event_hash` to compare against (TD-04).
 - **A reviewer's conflict of interest is not checked by the authorisation layer**, and submitted reviews are counted rather than identity-checked against an assignment (TD-02, TD-03, section 12.2) — both critical, both unresolved as of this document.
 - **Deployment runs on AWS root credentials.** Mitigated by never storing them as a CI secret, but not resolved (TD-01, section 12.2, section 13.4) — this is the single highest-priority open item in the project.
 - **The deployed architecture is smaller than the specified one.** App Runner, not ECS Fargate behind an ALB and CloudFront; no Redis; no asynchronous worker (TD-14, section 13.3). No functional capability is lost on the read/write paths that exist, but the asynchronous processing pipeline described in the design specification (section 10.2 of that document, and section 16 of this one) cannot run without the missing worker component.
 - **No load testing.** NFR-08 and NFR-09's performance objectives (archive pages within 500 ms p95, search within 800 ms p95) were sized at design time and have never been measured against the running system (section 11.6).
 - **No automated security scanning.** No SAST tool, dependency vulnerability scanner, or DAST pass runs in CI, despite both being named in the design specification (section 11) as intended gates (section 11.6, section 15.4).
 - **No browser-matrix or committed end-to-end testing.** UAT was run in a single Chromium-based session; `@playwright/test` is a declared dependency with no committed spec files (section 11.6).
 - **No file upload exists in the delivered domain.** `POST /manuscripts` is JSON-only — submission of an actual PDF, and the anonymisation/similarity pipeline that would process one, are specified (design specification section 10.2) but not built (`docs/05-api-contract.md` section 8, section 16 of this document).
 - **Explicit, permanent out-of-scope items** (design specification section 4.2), restated as hard limits rather than soft gaps: no payment or article-processing-charge handling; no copy-editing or typesetting workflow; no multi-journal tenancy (though the seam for it exists, section 16); no ORCID federation; identifiers are DOI-*shaped* but not registered with Crossref; similarity screening, when built, will be against the internal corpus only, never the open web; email deliverability is guaranteed through a single transactional provider only — and today, no live provider is wired in at all (`infrastructure/email/logging_sender.py` logs rather than sends, section 10.2).
-

18. Conclusion

18.1 What was estimated, and what was delivered

Use Case Points sized the Must-have subset of this system (UC1–UC18) at 2,514 person-hours, and the full specified system at 3,262 person-hours; COCOMO II's independent cross-check placed the full system at approximately 7,170 person-hours (section 7.2). The 48-hour window available for this final project is roughly 1.5–1.9% of either figure. The delivery — a working domain layer with 812 lines of framework-free, fully tested code; a complete API and frontend; a deployed, HTTPS-reachable system; and six supporting documents including this one — took a small fraction of the classically estimated effort. Section 7.5 explains why this does not invalidate the method: UCP measured the *problem's* size correctly, and what changed was the *rate* at which AI-assisted development converts that size into elapsed hours — a change in method, not evidence the sizing was wrong. The formal accounting of realised productivity is left to the mechanical procedure specified in effort estimation section 10.3, to be computed from commit history rather than asserted here.

18.2 What that trade bought, and what it cost

It bought a system that meets its three sharpest objectives (section 3) more convincingly than most final projects reach in the available time: **O1** (no illegal transition reachable through any interface) is enforced by an exhaustively tested state machine (section 8.2, section 11.2); **O2** (structural double-blind integrity) is enforced by a projection type with nowhere to put an author field (section 8.4); **O3** (a tamper-evident audit trail) is enforced by a hash chain and two database triggers, the second one added specifically because the first didn't cover every statement class (section 8.5, section 11.4). It bought a hexagonal architecture whose separation is not aspirational but mechanically checked on every commit by two import-linter contracts (section 9.2). And it bought a technical debt register and testing report that state their own limitations more precisely than most fully-staffed projects manage, because — per section 12.6 — this project's clearest lesson is about what automated rigour does and does not catch.

It cost what section 7.5 named in advance rather than discovered afterward: test depth below what a fully-priced 2,514-hour effort would buy (partially offset by the mutation-testing pass that found what coverage alone missed, section 11.4); documentation formality reduced to what six documents and inline code comments provide, rather than a fuller set of architecture decision records; and security hardening that verifies the NFR-01–NFR-06 baseline without attempting threat modelling or penetration testing beyond it. It cost the three critical, currently-open items in section 12.2 — root AWS credentials, an unenforced reviewer conflict-of-interest check, and an uncounted review quorum — none of which are hypothetical risks; all three are load-bearing gaps in a system whose stated purpose is enforcing exactly the guarantees they undermine.

18.3 The honest summary

The SDJ Editorial Portal demonstrates that a domain-first, hexagonally-architected system, built under AI-assisted development, can deliver mechanically-enforced correctness for the properties that were designed in from the start — the lifecycle, the blinding, the audit chain's internal consistency — within a timeframe that would be absurd for the same guarantees built by conventional estimation. It demonstrates equally clearly, through its own technical debt register and testing report, that passing every automated gate available to this project — linting, strict typing, an architecture contract, and a full test suite at 100% coverage — was necessary and not sufficient, and that the defects and gaps that mattered most were found by a human, or an agent acting as one, reading code against what it claimed to do, or by using the deployed system as an actual actor would. That finding, not the deployment URL, is this final project's most transferable result.

19. References

- IEEE Std 830-1998, *Recommended Practice for Software Requirements Specifications*.
- ISO/IEC/IEEE 29148:2018, *Systems and software engineering — Life cycle processes — Requirements engineering*.

- ISO/IEC 14764, *Software Engineering — Software Life Cycle Processes — Maintenance* (maintenance category taxonomy used in section 15).
- Karner, G. (1993). *Use Case Points* method for effort estimation, as codified in Cockburn, A. (2000), *Writing Effective Use Cases*, and Schneider & Winters (1998), *Applying Use Cases: A Practical Guide*.
- Boehm, B. et al. (2000). *Software Cost Estimation with COCOMO II* — Early Design model.
- Fowler, M. — technical debt quadrant (deliberate/inadvertent × reckless/prudent), the classification basis for `docs/04-technical-debt-register.md`.
- IETF RFC 9457, *Problem Details for HTTP APIs* — the platform's error-response format (section 9.3).
- Open Archives Initiative, *OAI-PMH 2.0 specification* (section 16, future evolution).

Project documents (this repository):

- `docs/02-srs.md` — Software Requirements Specification
- `docs/03-effort-estimation.md` — Effort estimation
- `docs/04-technical-debt-register.md` — Technical debt register
- `docs/05-api-contract.md` — API contract
- `docs/06-testing-report.md` — Testing report
- `docs/07-user-manual.md` — User manual (in progress at the time of writing; section 14 stands in until it is complete)
- `docs/superpowers/specs/2026-08-12-ugjcs-journal-platform-design.md` — Design specification

AI-assisted development acknowledgement. Consistent with the acknowledgement recorded in `docs/03-effort-estimation.md` section 9 and section 11, this document — and the code, tests and documents it consolidates — were produced with Claude Code (Anthropic), an AI coding assistant, under the direction and review of the author. The assistant drafted prose and code from the author's instructions and the design specification; the author directed the work, reviewed every output, and accepts sole responsibility for its correctness and for the claims and conclusions recorded in this document.